

# *Grafana & Flux*

## *New Flux support in Grafana*


*Jacob Lisi*  
*@JacobLisi*

# TL;DR

- Flux is powerful
- You can start playing with Flux in Grafana today
  - Flux support in Grafana is available via a new datasource plugin
- For now, no automated way to migrate dashboards
  - Transpiler for Influx queries is being worked on
  - You can migrate your dashboards and panels manually

# Flux Design Goals

- Decouple language from the execution engine
  - Execution engine takes a DAG of transformations
  - Transpilers: Flux, Influx Query, PromQL
- Decouple database from query computation
  - Iterate faster on query engine
  - Keep database untouched
  - Independently scalable
- Add more functions
  - Chainable: transformation from input table to output table



## *Flux recap*

```
from(db: "telegraf")  
  |> filter(fn: (r) => r["host"] == "myServer")  
  |> range(start: -1h)
```

## *Flux recap: Select one value from a table*

```
from(db: "telegraf")  
  |> filter(fn: (r) => r["host"] == "myServer")  
  |> range(start: -1h)  
  |> max() // Selector
```

## *Flux primer: Window-chunk and aggregate*

```
from(db: "telegraf")  
  |> filter(fn: (r) => r["_measurement"] == "cpu")  
  |> range(start: -1h)  
  |> window(every: 10m)  
  |> mean() // Aggregator  
  |> filter(fn: (r) => r._value > 1) // Having
```

# *Flux query planning*

```
from(db: "telegraf")  
  |> filter(fn: (r) => r["host"] == "myServer")  
  |> range(start: -1h)  
  |> max()
```

```
from(db: "telegraf")  
  |> range(start: -1h)  
  |> filter(fn: (r) => r["host"] == "myServer")  
  |> max()
```

- **Same plan DAG**

# *Flux query planning gotchas*

```
from(db: "telegraf")  
  |> filter(fn: (r) => r["host"] == "myServer")  
  |> range(start: -1h)  
  |> max()
```

```
from(db: "telegraf")  
  |> filter(fn: (r) => r["host"] == "myServer")  
  |> max() // Selector function returns 1 record  
  |> range(start: -1h)
```

```
from(db: "telegraf")  
  |> range(start: -1h)  
  |> filter(fn: (r) => r["_value"] > 1) // Full table scan
```



## *User defined functions*

```
select = (db="telegraf", m, f) => {  
  return from(db:db)  
    |> filter(fn: (r) => r._measurement == m  
      and r._field == f)  
}
```

```
select(m: "cpu", f: "usage_user")  
  |> filter(fn: (r) => r["host"] == "myServer")  
  |> range(start: -1h)
```

## *Chainable user defined functions*

```
myFilter = (m, f, table=<-) => {  
  return table  
    |> filter(fn: (r) => r._measurement == m  
      and r._field == f)  
}
```

```
from(db: "telegraf")  
  |> myFilter(m: "cpu", f: "usage_user")  
  |> range(start: -1h)
```

# Flux StdLib

<https://github.com/influxdata/flux/tree/master/stdlib>

The screenshot shows the GitHub interface for the `influxdata/flux` repository. The browser address bar displays the URL `https://github.com/influxdata/flux/blob/master/stdlib/universe/universe.flux`. The repository page header includes the `influxdata / flux` breadcrumb, navigation links for `Code`, `Issues` (315), `Pull requests` (12), and `Insights`. The file path `flux / stdlib / universe / universe.flux` is shown, along with buttons for `Find file` and `Copy path`. A commit by `aanthony1243` is highlighted, with the message `feat(stdlib/universe): add contains function to check for membership ...` and a commit hash of `0f58806` from 19 days ago. Below the commit, it lists `4 contributors`. The file details show `244 lines (217 sloc)` and `8.61 KB`. The file content is displayed in a code editor with tabs for `View file`, `Raw`, `Blame`, and `History`. The code is a Flux script defining a package, imports, and built-in functions.

```
1 package universe
2
3 import "system"
4
5 // now is a function option whose default behaviour is to return the current system time
6 option now = system.time
7
8 // Booleans
9 builtin true
10 builtin false
11
12 // Transformation functions
13 builtin columns
```

# *Math on tables*

```
cpu = from(db)...
```

```
CpuRequests = from(db)...
```

```
join(  
  tables: {cpu: cpu, req: CpuRequests},  
  fn: (t) => t.cpu._value / t.req._value  
) // Implicit join on time
```

# New response format: CSV

x		Headers	Preview	Response	Cookies	Timing
1		#datatype,string,long,dateTime:RFC3339,dateTime:RFC3339,dateTime:RFC3339,double,string,string,string,string				
2		#partition,false,false,true,true,false,false,true,true,true,true				
3		#default,result,,,,,,,,				
4		,result,table,start,stop,time,value,field,measurement,cpu,host				
5	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:46:15Z,95.35696455317024,usage_idle,cpu,cpu-total,kenobi				
6	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:46:25Z,95.04628471353516,usage_idle,cpu,cpu-total,kenobi				
7	...	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:46:35Z,95.15,usage_idle,cpu,cpu-total,kenobi				
8	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:46:45Z,95.27381845461365,usage_idle,cpu,cpu-total,kenobi				
9	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:46:55Z,95.10244877561219,usage_idle,cpu,cpu-total,kenobi				
10	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:47:05Z,94.2177722152691,usage_idle,cpu,cpu-total,kenobi				
11	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:47:15Z,55.563872255489024,usage_idle,cpu,cpu-total,kenobi				
12	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:47:25Z,89.68194340095167,usage_idle,cpu,cpu-total,kenobi				
13	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:47:35Z,94.20289855072464,usage_idle,cpu,cpu-total,kenobi				
14	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:47:45Z,94.57771114442778,usage_idle,cpu,cpu-total,kenobi				
15	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:47:55Z,92.04204204204204,usage_idle,cpu,cpu-total,kenobi				
16	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:48:05Z,93.975,usage_idle,cpu,cpu-total,kenobi				
17	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:48:15Z,94.81037924151697,usage_idle,cpu,cpu-total,kenobi				
18	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:48:25Z,94.56276622400401,usage_idle,cpu,cpu-total,kenobi				
19	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:48:35Z,95.43299226353881,usage_idle,cpu,cpu-total,kenobi				
20	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:48:45Z,95.26908635794743,usage_idle,cpu,cpu-total,kenobi				
21	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:48:55Z,95.25948103792415,usage_idle,cpu,cpu-total,kenobi				
22	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:49:05Z,95.54443053817272,usage_idle,cpu,cpu-total,kenobi				
23	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:49:15Z,95.49549549549549,usage_idle,cpu,cpu-total,kenobi				
24	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:49:25Z,95.45454545454545,usage_idle,cpu,cpu-total,kenobi				
25	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:49:35Z,95.3953953953954,usage_idle,cpu,cpu-total,kenobi				
26	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:49:45Z,95.30352235823133,usage_idle,cpu,cpu-total,kenobi				
27	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:49:55Z,95.37615596100974,usage_idle,cpu,cpu-total,kenobi				
28	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:50:05Z,95.15726410384423,usage_idle,cpu,cpu-total,kenobi				
29	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:50:15Z,94.7908840470824,usage_idle,cpu,cpu-total,kenobi				
30	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:50:25Z,95.39769884942471,usage_idle,cpu,cpu-total,kenobi				
31	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:50:35Z,95.03121098626717,usage_idle,cpu,cpu-total,kenobi				
32	..	,,0,2018-06-13T12:46:14.295154162Z,2018-06-13T18:46:14.295154162Z,2018-06-13T12:50:45Z,95.15726410384423,usage_idle,cpu,cpu-total,kenobi				

# Getting Started With Flux

- Run latest influxd
  - <https://portal.influxdata.com/downloads>
  - Update your influxdb.conf to include

```
# ...  
[http]  
# ...  
flux-enabled = true  
# ...
```

- Generate data
  - Telegraf

# *Get started: Grafana datasource*

- Get Grafana 5.3+
- Install Flux datasource plugin
  - <https://github.com/grafana/influxdb-flux-datasource>
  - Clone into your grafanas data/plugins
  - Restart Grafana
- Add your Flux datasource
- Add a dashboard
- Add a panel

# *Demo*

[https://github.com/jtlisi/grafana\\_flux\\_demo](https://github.com/jtlisi/grafana_flux_demo)



# Datasource feature summary

- Syntax highlighting, tab completion, raw table preview
- Inline function documentation
- \$range variable

The screenshot displays the Grafana Labs interface with the 'Table' view selected. The 'Metrics' tab is active, and the 'Data Source' is set to 'Flux'. The query editor shows a query with syntax highlighting and a dropdown menu for the 'map' function. The dropdown menu lists the function 'map(fn: (r) => r)' with its description 'Applies a function to each record of the input tables.' and the 'max()' function. The 'Result records' are displayed as 5000.

```
from(db: "telegraf")
|> filter(fn: (r) => r["_measurement"] == "cpu" AND
r["_field"] == "usage_idle")
|> range($range)
|> limit(n: 1000)
|> ma
```

Functions

- map(fn: (r) => r)  
Applies a function to each record of the input tables.
- max()

Result records 5000

# Datasource feature summary

- Shortcodes
- Template variables with helper functions
  - measurements()
  - field\_keys()
  - tags()
  - tag\_values()
- Annotations

Variables > Edit

General

Name	Hosts	Type	Query
Label	optional display name	Hide	

Query Options

Data source	Flux	Refresh	Never
Query	tag_values(telegraf,cpu,host)		

# *Roadmap*

- Alerting
- UI improvements
- Improve Query Shortcuts
- Transitioning to a default plugin
- Dashboard Migrations?

*Thanks for listening!*  
*Questions?*

@JacobLisi